

Interviewly — scaling, latency and performance report

run window	2026-08-12T17:14:59.327Z to 2026-08-12T17:25:08.158Z
git sha	a0a167632cdb90d50956d00fd43975538125775f
target	docker-compose
base url	http://localhost
source file	/Users/sezaiemrekonuk/Desktop/obss/project/loadtest/results/compose-scale.json
schema	interviewly-loadtest/1
runs recorded	36

1. Environment

property	value
host platform	darwin 25.5.0
host arch	arm64
host cpus	8
host memory	16.0 GiB
node	25.9.0
docker cpus	8
docker memory	7.8 GiB
docker server version	29.4.0
docker operating system	Docker Desktop

Caveat: the load generator and the system under test share this one machine. Generator CPU competes with api, web, database and cache containers for the same cores, so absolute throughput is a floor, not a ceiling, and latency includes contention the generator itself creates.

2. Method

Scenarios exercised

scenario	method	url
healthz	GET	http://localhost/api/healthz
readyz	GET	http://localhost/api/readyz
me	GET	http://localhost/api/me
my-interviews	GET	http://localhost/api/me/interviews
interview-state	GET	http://localhost/api/interviews/cmsqcavkv0001n001h4jtx9pc/state
web-home	GET	http://localhost/

Generator configuration

duration per cell	12,000 ms
warmup per cell	3,000 ms
connection levels	8, 64
replica counts	1, 2, 4
generator	node fetch, closed loop, one in-flight request per connection
ai enabled	false
log transport	stdout
db pool per process	10
compose files	compose.yaml, compose.scale.yaml
interview id	cmsqcavkv0001n001h4jtx9pc

3. Instance fan-out proof

Each scale step scaled the api service, then probed the edge until every replica had answered. The counts below are the evidence that N replicas actually served traffic rather than one replica absorbing it all.

replicas requested	containers	distinct instances via edge	instance ids	at
1	1	1	c2f71fa0	2026-08-12T17:15:11.72Z
2	2	2	c2f71fa0, 6a5c890c	2026-08-12T17:18:36.981Z
4	4	4	c7419cf9, 6a5c890c, c2f71fa0, a70da523	2026-08-12T17:22:03.201Z

Per-run share of requests by instance

replicas	scenario	conns	requests	instance share
1	healthz	8	75,987	c2f71fa0 100.0% (75,987)
1	healthz	64	74,207	c2f71fa0 100.0% (74,207)
1	readyz	8	52,382	c2f71fa0 100.0% (52,382)
1	readyz	64	54,626	c2f71fa0 100.0% (54,626)
1	me	8	16,444	c2f71fa0 100.0% (16,444)
1	me	64	13,172	c2f71fa0 100.0% (13,172)
1	my-interviews	8	4,608	c2f71fa0 100.0% (4,608)
1	my-interviews	64	4,686	c2f71fa0 100.0% (4,686)
1	interview-state	8	2,001	c2f71fa0 100.0% (2,001)
1	interview-state	64	1,709	c2f71fa0 100.0% (1,709)
1	web-home	8	1,459	none 100.0% (1,459)
1	web-home	64	1,508	none 100.0% (1,508)
2	healthz	8	75,844	c2f71fa0 50.0% (37,923) · 6a5c890c 50.0% (37,921)
2	healthz	64	68,438	6a5c890c 50.0% (34,243) · c2f71fa0 50.0% (34,195)
2	readyz	8	40,348	c2f71fa0 50.0% (20,175) · 6a5c890c 50.0% (20,173)
2	readyz	64	42,027	6a5c890c 50.0% (21,018) · c2f71fa0 50.0% (21,009)
2	me	8	14,537	c2f71fa0 50.0% (7,269) · 6a5c890c 50.0% (7,268)
2	me	64	14,226	6a5c890c 50.2% (7,136) · c2f71fa0 49.8% (7,090)
2	my-interviews	8	6,250	c2f71fa0 50.0% (3,127) · 6a5c890c 50.0% (3,123)
2	my-interviews	64	6,163	6a5c890c 50.4% (3,108) · c2f71fa0 49.6% (3,055)
2	interview-state	8	2,694	c2f71fa0 50.0% (1,348) · 6a5c890c 50.0% (1,346)
2	interview-state	64	2,896	6a5c890c 50.5% (1,463) · c2f71fa0 49.5% (1,433)
2	web-home	8	1,250	none 100.0% (1,250)
2	web-home	64	1,522	none 100.0% (1,522)
4	healthz	8	71,734	a70da523 25.0% (17,934) · c7419cf9 25.0% (17,934) · c2f71fa0 25.0% (17,933) · 6a5c890c 25.0% (17,933)
4	healthz	64	84,575	c7419cf9 25.0% (21,158) · 6a5c890c 25.0% (21,142) · c2f71fa0 25.0% (21,139) · a70da523 25.0% (21,136)
4	readyz	8	36,748	c2f71fa0 25.0% (9,188) · 6a5c890c 25.0% (9,188) · c7419cf9 25.0% (9,186) · a70da523 25.0% (9,186)
4	readyz	64	48,945	6a5c890c 25.0% (12,240) · a70da523 25.0% (12,237) · c2f71fa0 25.0% (12,235) · c7419cf9 25.0% (12,233)
4	me	8	14,291	c7419cf9 25.0% (3,573) · c2f71fa0 25.0% (3,573) · a70da523 25.0% (3,573) · 6a5c890c 25.0% (3,572)
4	me	64	17,174	c7419cf9 25.1% (4,319) · 6a5c890c 25.0% (4,288) · a70da523 25.0% (4,285) · c2f71fa0 24.9% (4,282)
4	my-interviews	8	7,694	a70da523 25.0% (1,924) · 6a5c890c 25.0% (1,924) · c7419cf9 25.0% (1,923) · c2f71fa0 25.0% (1,923)
4	my-interviews	64	5,345	6a5c890c 25.2% (1,346) · c7419cf9 25.0% (1,337) · a70da523 24.9% (1,332) · c2f71fa0 24.9% (1,330)
4	interview-state	8	2,568	c2f71fa0 25.0% (643) · 6a5c890c 25.0% (642) · a70da523 25.0% (642) · c7419cf9 25.0% (641)
4	interview-state	64	3,608	6a5c890c 25.4% (915) · c2f71fa0 24.9% (898) · a70da523 24.9% (898) · c7419cf9 24.9% (897)

replicas	scenario	conns	requests	instance share
4	web-home	8	1,446	none 100.0% (1,446)
4	web-home	64	1,383	none 100.0% (1,383)

4. Throughput scaling

Requests per second — connections = 8

scenario	1 replicas	2 replicas	4 replicas	factor 4x / 1x
healthz	6580	6540	6239	0.95
readyz	4473	3482	3182	0.71
me	1415	1256	1254	0.89
my-interviews	393	540	674	1.71
interview-state	171	232	231	1.35
web-home	125	107	125	0.99

Requests per second — connections = 64

scenario	1 replicas	2 replicas	4 replicas	factor 4x / 1x
healthz	6400	5974	7462	1.17
readyz	4698	3685	4437	0.94
me	1129	1247	1518	1.34
my-interviews	400	535	524	1.31
interview-state	144	248	324	2.24
web-home	128	128	117	0.92

5. Latency

Client latency is measured by the generator around the whole request. Server latency comes from the in-process route profiler; where several replicas served a scenario the worst instance is reported. The overhead column is client p95 minus server p95 on that worst instance.

healthz

replicas	conns	client p50	client p95	client p99	client max	busiest route	srv p50	srv p95	srv p99	edge+net p95
1	8	1.02	2.47	4.47	13.50	GET /healthz	0.08	0.17	0.29	2.30
2	8	0.97	2.64	4.96	36.07	GET /healthz	0.13	0.39	0.67	2.25
4	8	1.00	2.81	5.61	26.14	GET /healthz	0.16	0.47	0.77	2.34
1	64	8.75	16.32	31.15	557.43	GET /healthz	0.08	0.19	0.32	16.13
2	64	7.06	22.72	115.73	690.78	GET /healthz	0.15	0.40	0.68	22.32
4	64	5.73	22.79	47.18	266.73	GET /healthz	0.23	0.76	1.99	22.03

readyz

replicas	conns	client p50	client p95	client p99	client max	busiest route	srv p50	srv p95	srv p99	edge+net p95
1	8	1.59	2.96	5.49	16.44	GET /readyz	0.55	1.09	1.95	1.87
2	8	1.97	4.26	7.51	29.49	GET /readyz	0.89	2.24	4.39	2.02
4	8	2.12	4.91	8.43	42.71	GET /readyz	1.03	2.70	5.15	2.20
1	64	12.41	20.57	31.48	224.87	GET /readyz	4.37	7.91	11.29	12.65
2	64	14.23	36.02	94.30	323.75	GET /readyz	4.43	11.95	23.42	24.07
4	64	11.51	33.09	59.29	273.78	GET /readyz	3.79	13.94	24.71	19.15

me

replicas	conns	client p50	client p95	client p99	client max	busiest route	srv p50	srv p95	srv p99	edge+net p95
1	8	5.38	8.42	12.40	31.34	GET /me	3.88	5.71	6.72	2.71
2	8	5.79	10.76	16.41	84.13	GET /me	5.44	9.51	12.71	1.26
4	8	5.60	11.43	18.09	47.58	GET /me	4.32	9.08	14.92	2.36
1	64	53.75	75.75	97.46	240.24	GET /me	30.42	42.22	50.00	33.53
2	64	57.66	102.77	124.85	159.83	GET /me	52.55	73.13	89.60	29.64
4	64	32.02	100.08	157.28	497.90	GET /me	27.50	71.32	101.59	28.75

my-interviews

replicas	conns	client p50	client p95	client p99	client max	busiest route	srv p50	srv p95	srv p99	edge+net p95
1	8	20.32	27.59	35.18	107.05	GET /me/interviews	17.88	23.98	30.89	3.61
2	8	14.15	24.51	30.51	60.72	GET /me/interviews	15.44	23.33	27.94	1.18
4	8	11.01	18.64	26.23	69.97	GET /me/interviews	9.36	16.69	25.84	1.95
1	64	156.52	197.18	222.50	271.46	GET /me/interviews	135.23	168.02	188.08	29.17
2	64	135.20	231.57	286.96	338.53	GET /me/interviews	132.38	186.85	230.18	44.72
4	64	120.01	225.67	272.81	353.31	GET /me/interviews	108.70	218.46	282.29	7.21

interview-state


replicas	conns	client p50	client p95	client p99	client max	busiest route	srv p50	srv p95	srv p99	edge+net p95
1	8	45.39	59.08	85.82	117.12	GET /interviews/:id/state	42.58	56.08	79.05	3.00
2	8	34.03	47.45	56.99	79.16	GET /interviews/:id/state	35.68	47.58	58.01	-0.13
4	8	31.53	57.60	86.23	136.78	GET /interviews/:id/state	30.52	58.30	91.93	-0.70
1	64	421.88	611.43	949.91	990.45	GET /interviews/:id/state	396.20	566.66	872.89	44.77
2	64	307.67	503.42	571.43	618.97	GET /interviews/:id/state	357.97	480.11	522.18	23.31
4	64	120.85	530.29	586.44	660.63	GET /interviews/:id/state	422.65	514.90	560.62	15.38

web-home


replicas	conns	client p50	client p95	client p99	client max	busiest route	srv p50	srv p95	srv p99	edge+net p95
1	8	60.06	91.86	132.18	140.03	—	—	—	—	—
2	8	65.96	127.40	264.81	463.07	—	—	—	—	—
4	8	62.00	87.31	111.14	186.65	—	—	—	—	—
1	64	497.17	670.52	854.43	911.41	—	—	—	—	—
2	64	462.34	954.77	1390.74	1607.53	—	—	—	—	—
4	64	536.97	879.01	960.29	1172.84	—	—	—	—	—

6. Charts

Throughput by replica count — connections = 64



Client p95 latency by replica count — connections = 64



7. Resource use

Sampled mid-run at the highest connection level (connections = 64). Docker CPU percentages are per container as reported by docker stats; the sum column adds the api containers together.

scenario	replicas	api cpu sum	event-loop p99 per instance	pg conns	redis clients	xact commits
healthz	1	112.0%	c2f71fa0 26.56 ms	12	9	14
readyz	1	159.9%	c2f71fa0 20.51 ms	12	9	243
me	1	444.4%	c2f71fa0 28.34 ms	12	9	37,986
my-interviews	1	572.7%	c2f71fa0 48.53 ms	12	9	18,927
interview-state	1	572.5%	c2f71fa0 91.55 ms	12	9	29,285
web-home	1	0.3%	c2f71fa0 13.43 ms	12	9	13
healthz	2	178.8%	6a5c890c 35.62 ms · c2f71fa0 29.84 ms	13	11	20
readyz	2	304.6%	c2f71fa0 33.37 ms · 6a5c890c 31.70 ms	22	11	364
me	2	500.6%	c2f71fa0 16.14 ms · 6a5c890c 40.40 ms	22	11	41,348
my-interviews	2	561.3%	c2f71fa0 26.98 ms · 6a5c890c 87.95 ms	22	11	24,319
interview-state	2	599.5%	6a5c890c 73.53 ms · c2f71fa0 27.56 ms	22	11	49,814

scenario	replicas	api cpu sum	event-loop p99 per instance	pg conns	redis clients	xact commits
web-home	2	0.5%	c2f71fa0 15.68 ms · 6a5c890c 15.42 ms	22	11	14
healthz	4	333.2%	6a5c890c 24.58 ms · c2f71fa0 26.49 ms · a70da523 26.48 ms · c7419cf9 27.39 ms	24	15	27
readyz	4	375.5%	6a5c890c 27.56 ms · c7419cf9 29.11 ms · c2f71fa0 24.35 ms · a70da523 25.72 ms	42	15	1,573
me	4	503.8%	c2f71fa0 29.90 ms · 6a5c890c 41.16 ms · c7419cf9 41.75 ms · a70da523 30.80 ms	42	15	51,600
my-interviews	4	539.7%	6a5c890c 86.11 ms · c7419cf9 89.85 ms · c2f71fa0 74.06 ms · a70da523 75.17 ms	42	15	21,482
interview-state	4	545.7%	6a5c890c 81.53 ms · c7419cf9 32.01 ms · a70da523 32.90 ms · c2f71fa0 34.44 ms	42	15	61,761
web-home	4	5.9%	c7419cf9 14.89 ms · a70da523 13.92 ms · c2f71fa0 14.53 ms · 6a5c890c 14.60 ms	42	15	14

Container detail

scenario	replicas	container	cpu	memory	mem %	net io
healthz	1	interviewly-worker-1	0.1%	51.06MiB / 7.75GiB	0.6%	232kB / 513kB
healthz	1	interviewly-api-1	112.0%	186.4MiB / 7.75GiB	2.4%	2.13GB / 751MB
healthz	1	interviewly-bucket-1	0.2%	90.86MiB / 7.75GiB	1.1%	43.1kB / 10.8kB
healthz	1	interviewly-mail-1	0.0%	7.727MiB / 7.75GiB	0.1%	9.7kB / 126B
healthz	1	interviewly-web-1	0.0%	66MiB / 7.75GiB	0.8%	7.55MB / 247MB
healthz	1	interviewly-kibana-1	5.9%	594.4MiB / 7.75GiB	7.5%	14.6MB / 87.4MB
healthz	1	interviewly-elasticsearch-1	1.0%	1.043GiB / 7.75GiB	13.5%	83.2MB / 14MB
healthz	1	interviewly-db-1	0.0%	65.04MiB / 7.75GiB	0.8%	181MB / 2.27GB
healthz	1	interviewly-cache-1	0.5%	13.65MiB / 7.75GiB	0.2%	43.8MB / 32.1MB
healthz	1	interviewly-edge-1	125.0%	35.25MiB / 7.75GiB	0.4%	1.37GB / 1.53GB
readyz	1	interviewly-worker-1	0.0%	51.06MiB / 7.75GiB	0.6%	236kB / 526kB
readyz	1	interviewly-api-1	159.9%	187.7MiB / 7.75GiB	2.4%	2.2GB / 822MB
readyz	1	interviewly-bucket-1	0.1%	90.86MiB / 7.75GiB	1.1%	43.1kB / 10.8kB
readyz	1	interviewly-mail-1	0.0%	7.727MiB / 7.75GiB	0.1%	9.7kB / 126B
readyz	1	interviewly-web-1	0.0%	67.77MiB / 7.75GiB	0.8%	7.55MB / 247MB
readyz	1	interviewly-kibana-1	2.8%	594.4MiB / 7.75GiB	7.5%	14.6MB / 87.6MB
readyz	1	interviewly-elasticsearch-1	0.7%	1.043GiB / 7.75GiB	13.5%	83.5MB / 14.1MB
readyz	1	interviewly-db-1	13.5%	58.32MiB / 7.75GiB	0.7%	193MB / 2.29GB
readyz	1	interviewly-cache-1	9.0%	13.66MiB / 7.75GiB	0.2%	53MB / 40.5MB
readyz	1	interviewly-edge-1	89.2%	34.92MiB / 7.75GiB	0.4%	1.46GB / 1.64GB
me	1	interviewly-worker-1	0.1%	51.07MiB / 7.75GiB	0.6%	240kB / 541kB
me	1	interviewly-api-1	444.4%	192.3MiB / 7.75GiB	2.4%	2.37GB / 870MB
me	1	interviewly-bucket-1	0.0%	90.79MiB / 7.75GiB	1.1%	43.1kB / 10.8kB
me	1	interviewly-mail-1	0.0%	7.727MiB / 7.75GiB	0.1%	9.7kB / 126B
me	1	interviewly-web-1	0.0%	69.27MiB / 7.75GiB	0.9%	7.55MB / 247MB
me	1	interviewly-kibana-1	2.5%	594.4MiB / 7.75GiB	7.5%	14.6MB / 87.9MB
me	1	interviewly-elasticsearch-1	0.9%	1.043GiB / 7.75GiB	13.5%	83.7MB / 14.1MB
me	1	interviewly-db-1	25.1%	59.19MiB / 7.75GiB	0.8%	214MB / 2.43GB
me	1	interviewly-cache-1	0.4%	13.65MiB / 7.75GiB	0.2%	54.5MB / 41.8MB
me	1	interviewly-edge-1	27.3%	33.57MiB / 7.75GiB	0.4%	1.5GB / 1.69GB
my-interviews	1	interviewly-worker-1	0.5%	51.06MiB / 7.75GiB	0.6%	244kB / 556kB
my-interviews	1	interviewly-api-1	572.7%	206.4MiB / 7.75GiB	2.6%	2.85GB / 933MB
my-interviews	1	interviewly-bucket-1	0.0%	90.79MiB / 7.75GiB	1.1%	43.1kB / 10.8kB
my-interviews	1	interviewly-mail-1	0.0%	7.734MiB / 7.75GiB	0.1%	9.7kB / 126B

scenario	replicas	container	cpu	memory	mem %	net io
my-interviews	1	interviewly-web-1	0.0%	68.7MiB / 7.75GiB	0.9%	7.55MB / 247MB
my-interviews	1	interviewly-kibana-1	3.3%	594.4MiB / 7.75GiB	7.5%	14.7MB / 88.1MB
my-interviews	1	interviewly-elasticsearch-1	1.7%	1.043GiB / 7.75GiB	13.5%	84MB / 14.1MB
my-interviews	1	interviewly-db-1	35.1%	60.55MiB / 7.75GiB	0.8%	230MB / 2.91GB
my-interviews	1	interviewly-cache-1	1.4%	13.65MiB / 7.75GiB	0.2%	54.5MB / 41.8MB
my-interviews	1	interviewly-edge-1	14.5%	30.98MiB / 7.75GiB	0.4%	1.55GB / 1.75GB
interview-state	1	interviewly-worker-1	0.2%	51.06MiB / 7.75GiB	0.6%	248kB / 570kB
interview-state	1	interviewly-api-1	572.5%	210.4MiB / 7.75GiB	2.6%	3.08GB / 1.04GB
interview-state	1	interviewly-bucket-1	0.1%	90.81MiB / 7.75GiB	1.1%	43.1kB / 10.8kB
interview-state	1	interviewly-mail-1	0.0%	7.742MiB / 7.75GiB	0.1%	9.7kB / 126B
interview-state	1	interviewly-web-1	0.0%	68.94MiB / 7.75GiB	0.9%	7.55MB / 247MB
interview-state	1	interviewly-kibana-1	2.8%	594.4MiB / 7.75GiB	7.5%	14.7MB / 88.4MB
interview-state	1	interviewly-elasticsearch-1	0.8%	1.045GiB / 7.75GiB	13.5%	84.2MB / 14.2MB
interview-state	1	interviewly-db-1	35.1%	66.59MiB / 7.75GiB	0.8%	247MB / 3.14GB
interview-state	1	interviewly-cache-1	1.8%	13.65MiB / 7.75GiB	0.2%	55.1MB / 42.1MB
interview-state	1	interviewly-edge-1	7.9%	31.28MiB / 7.75GiB	0.4%	1.65GB / 1.84GB
web-home	1	interviewly-worker-1	1.5%	51.07MiB / 7.75GiB	0.6%	252kB / 584kB
web-home	1	interviewly-api-1	0.3%	65.18MiB / 7.75GiB	0.8%	3.11GB / 1.06GB
web-home	1	interviewly-bucket-1	0.0%	90.8MiB / 7.75GiB	1.1%	43.1kB / 10.8kB
web-home	1	interviewly-mail-1	0.0%	7.742MiB / 7.75GiB	0.1%	9.7kB / 126B
web-home	1	interviewly-web-1	148.4%	312.1MiB / 7.75GiB	3.9%	11.1MB / 366MB
web-home	1	interviewly-kibana-1	2.9%	594.4MiB / 7.75GiB	7.5%	14.7MB / 88.6MB
web-home	1	interviewly-elasticsearch-1	2.8%	1.045GiB / 7.75GiB	13.5%	84.5MB / 14.2MB
web-home	1	interviewly-db-1	0.0%	66.04MiB / 7.75GiB	0.8%	249MB / 3.17GB
web-home	1	interviewly-cache-1	0.5%	13.65MiB / 7.75GiB	0.2%	55.2MB / 42.1MB
web-home	1	interviewly-edge-1	15.4%	34.08MiB / 7.75GiB	0.4%	1.79GB / 1.98GB
healthz	2	interviewly-api-2	80.0%	151.7MiB / 7.75GiB	1.9%	26MB / 27.3MB
healthz	2	interviewly-worker-1	0.0%	51.07MiB / 7.75GiB	0.6%	258kB / 608kB
healthz	2	interviewly-api-1	98.8%	183.7MiB / 7.75GiB	2.3%	3.14GB / 1.09GB
healthz	2	interviewly-bucket-1	0.0%	90.83MiB / 7.75GiB	1.1%	43.5kB / 10.8kB
healthz	2	interviewly-mail-1	0.0%	7.742MiB / 7.75GiB	0.1%	10.1kB / 126B
healthz	2	interviewly-web-1	8.3%	64.54MiB / 7.75GiB	0.8%	11.8MB / 388MB
healthz	2	interviewly-kibana-1	5.8%	594.6MiB / 7.75GiB	7.5%	14.8MB / 89MB
healthz	2	interviewly-elasticsearch-1	1.8%	1.047GiB / 7.75GiB	13.5%	84.9MB / 14.3MB
healthz	2	interviewly-db-1	0.0%	67.85MiB / 7.75GiB	0.8%	249MB / 3.17GB
healthz	2	interviewly-cache-1	0.4%	13.68MiB / 7.75GiB	0.2%	55.2MB / 42.1MB
healthz	2	interviewly-edge-1	140.4%	34.81MiB / 7.75GiB	0.4%	1.92GB / 2.13GB
readyz	2	interviewly-api-2	149.0%	156MiB / 7.75GiB	2.0%	53.8MB / 56.4MB
readyz	2	interviewly-worker-1	0.0%	51.07MiB / 7.75GiB	0.6%	262kB / 622kB
readyz	2	interviewly-api-1	155.5%	187.4MiB / 7.75GiB	2.4%	3.16GB / 1.11GB
readyz	2	interviewly-bucket-1	0.1%	90.86MiB / 7.75GiB	1.1%	43.5kB / 10.8kB
readyz	2	interviewly-mail-1	0.0%	7.742MiB / 7.75GiB	0.1%	10.1kB / 126B
readyz	2	interviewly-web-1	0.0%	65.72MiB / 7.75GiB	0.8%	11.8MB / 388MB
readyz	2	interviewly-kibana-1	5.0%	594.6MiB / 7.75GiB	7.5%	14.8MB / 89.3MB
readyz	2	interviewly-elasticsearch-1	2.0%	1.047GiB / 7.75GiB	13.5%	85.1MB / 14.3MB
readyz	2	interviewly-db-1	25.8%	82.23MiB / 7.75GiB	1.0%	259MB / 3.18GB
readyz	2	interviewly-cache-1	14.9%	13.65MiB / 7.75GiB	0.2%	62.8MB / 48.8MB
readyz	2	interviewly-edge-1	138.8%	34.87MiB / 7.75GiB	0.4%	2GB / 2.23GB

scenario	replicas	container	cpu	memory	mem %	net io
me	2	interviewly-api-2	306.5%	161.6MiB / 7.75GiB	2.0%	135MB / 79.1MB
me	2	interviewly-worker-1	0.0%	51.07MiB / 7.75GiB	0.6%	266kB / 636kB
me	2	interviewly-api-1	194.1%	188MiB / 7.75GiB	2.4%	3.25GB / 1.14GB
me	2	interviewly-bucket-1	0.0%	90.86MiB / 7.75GiB	1.1%	43.5kB / 10.8kB
me	2	interviewly-mail-1	0.0%	7.75MiB / 7.75GiB	0.1%	10.1kB / 126B
me	2	interviewly-web-1	0.0%	67.09MiB / 7.75GiB	0.8%	11.8MB / 388MB
me	2	interviewly-kibana-1	4.3%	594.6MiB / 7.75GiB	7.5%	14.9MB / 89.5MB
me	2	interviewly-elasticsearch-1	0.7%	1.047GiB / 7.75GiB	13.5%	85.3MB / 14.3MB
me	2	interviewly-db-1	41.9%	85.48MiB / 7.75GiB	1.1%	279MB / 3.32GB
me	2	interviewly-cache-1	0.5%	13.64MiB / 7.75GiB	0.2%	64MB / 49.8MB
me	2	interviewly-edge-1	43.9%	32.54MiB / 7.75GiB	0.4%	2.04GB / 2.28GB
my-interviews	2	interviewly-api-2	341.1%	185.5MiB / 7.75GiB	2.3%	462MB / 121MB
my-interviews	2	interviewly-worker-1	0.0%	51.07MiB / 7.75GiB	0.6%	270kB / 650kB
my-interviews	2	interviewly-api-1	220.2%	194.5MiB / 7.75GiB	2.5%	3.57GB / 1.18GB
my-interviews	2	interviewly-bucket-1	0.0%	90.89MiB / 7.75GiB	1.1%	43.5kB / 10.8kB
my-interviews	2	interviewly-mail-1	0.0%	7.762MiB / 7.75GiB	0.1%	10.1kB / 126B
my-interviews	2	interviewly-web-1	0.0%	68.63MiB / 7.75GiB	0.9%	11.8MB / 388MB
my-interviews	2	interviewly-kibana-1	3.8%	594.6MiB / 7.75GiB	7.5%	14.9MB / 89.7MB
my-interviews	2	interviewly-elasticsearch-1	1.5%	1.047GiB / 7.75GiB	13.5%	85.6MB / 14.4MB
my-interviews	2	interviewly-db-1	55.4%	87.92MiB / 7.75GiB	1.1%	300MB / 3.97GB
my-interviews	2	interviewly-cache-1	0.4%	13.64MiB / 7.75GiB	0.2%	64MB / 49.8MB
my-interviews	2	interviewly-edge-1	21.1%	32.36MiB / 7.75GiB	0.4%	2.11GB / 2.35GB
interview-state	2	interviewly-api-2	335.2%	188MiB / 7.75GiB	2.4%	626MB / 199MB
interview-state	2	interviewly-worker-1	0.2%	51.13MiB / 7.75GiB	0.6%	276kB / 671kB
interview-state	2	interviewly-api-1	264.3%	194.1MiB / 7.75GiB	2.5%	3.74GB / 1.26GB
interview-state	2	interviewly-bucket-1	0.0%	90.89MiB / 7.75GiB	1.1%	43.5kB / 10.8kB
interview-state	2	interviewly-mail-1	0.0%	7.762MiB / 7.75GiB	0.1%	10.1kB / 126B
interview-state	2	interviewly-web-1	0.0%	69.16MiB / 7.75GiB	0.9%	11.8MB / 388MB
interview-state	2	interviewly-kibana-1	3.8%	594.9MiB / 7.75GiB	7.5%	15MB / 90MB
interview-state	2	interviewly-elasticsearch-1	2.5%	1.047GiB / 7.75GiB	13.5%	85.9MB / 14.4MB
interview-state	2	interviewly-db-1	65.2%	91.63MiB / 7.75GiB	1.1%	325MB / 4.29GB
interview-state	2	interviewly-cache-1	1.2%	13.66MiB / 7.75GiB	0.2%	64.9MB / 50.2MB
interview-state	2	interviewly-edge-1	13.2%	32.55MiB / 7.75GiB	0.4%	2.25GB / 2.5GB
web-home	2	interviewly-api-2	0.2%	43.24MiB / 7.75GiB	0.5%	649MB / 214MB
web-home	2	interviewly-worker-1	0.0%	51.14MiB / 7.75GiB	0.6%	280kB / 684kB
web-home	2	interviewly-api-1	0.3%	64.95MiB / 7.75GiB	0.8%	3.76GB / 1.27GB
web-home	2	interviewly-bucket-1	0.1%	90.85MiB / 7.75GiB	1.1%	43.5kB / 10.8kB
web-home	2	interviewly-mail-1	0.0%	7.77MiB / 7.75GiB	0.1%	10.1kB / 126B
web-home	2	interviewly-web-1	157.8%	287.2MiB / 7.75GiB	3.6%	15.2MB / 505MB
web-home	2	interviewly-kibana-1	1.9%	595MiB / 7.75GiB	7.5%	15MB / 90.3MB
web-home	2	interviewly-elasticsearch-1	1.4%	1.047GiB / 7.75GiB	13.5%	86.1MB / 14.5MB
web-home	2	interviewly-db-1	1.3%	91.69MiB / 7.75GiB	1.2%	329MB / 4.34GB
web-home	2	interviewly-cache-1	1.0%	13.65MiB / 7.75GiB	0.2%	65MB / 50.3MB
web-home	2	interviewly-edge-1	13.8%	35.18MiB / 7.75GiB	0.4%	2.39GB / 2.64GB
healthz	4	interviewly-api-4	80.0%	146.6MiB / 7.75GiB	1.9%	14.3MB / 15MB
healthz	4	interviewly-api-3	87.4%	152.1MiB / 7.75GiB	1.9%	14.3MB / 15MB
healthz	4	interviewly-api-2	83.7%	156.1MiB / 7.75GiB	2.0%	664MB / 229MB
healthz	4	interviewly-worker-1	0.1%	51.14MiB / 7.75GiB	0.6%	287kB / 708kB

scenario	replicas	container	cpu	memory	mem %	net io
healthz	4	interviewly-api-1	82.1%	180.1MiB / 7.75GiB	2.3%	3.77GB / 1.29GB
healthz	4	interviewly-bucket-1	0.0%	90.88MiB / 7.75GiB	1.1%	44.1kB / 10.8kB
healthz	4	interviewly-mail-1	0.0%	7.77MiB / 7.75GiB	0.1%	10.7kB / 126B
healthz	4	interviewly-web-1	0.0%	65.07MiB / 7.75GiB	0.8%	15.7MB / 521MB
healthz	4	interviewly-kibana-1	4.2%	595MiB / 7.75GiB	7.5%	15.1MB / 90.6MB
healthz	4	interviewly-elasticsearch-1	0.6%	1.047GiB / 7.75GiB	13.5%	86.5MB / 14.5MB
healthz	4	interviewly-db-1	0.0%	94.34MiB / 7.75GiB	1.2%	329MB / 4.34GB
healthz	4	interviewly-cache-1	0.8%	13.69MiB / 7.75GiB	0.2%	65.1MB / 50.3MB
healthz	4	interviewly-edge-1	205.5%	36.75MiB / 7.75GiB	0.5%	2.53GB / 2.8GB
readyz	4	interviewly-api-4	92.4%	152.8MiB / 7.75GiB	1.9%	28.5MB / 29.8MB
readyz	4	interviewly-api-3	92.6%	154.4MiB / 7.75GiB	1.9%	28.5MB / 29.8MB
readyz	4	interviewly-api-2	94.6%	159.6MiB / 7.75GiB	2.0%	678MB / 244MB
readyz	4	interviewly-worker-1	0.0%	51.15MiB / 7.75GiB	0.6%	291kB / 722kB
readyz	4	interviewly-api-1	95.9%	185.1MiB / 7.75GiB	2.3%	3.79GB / 1.3GB
readyz	4	interviewly-bucket-1	0.0%	90.87MiB / 7.75GiB	1.1%	44.1kB / 10.8kB
readyz	4	interviewly-mail-1	0.0%	7.773MiB / 7.75GiB	0.1%	10.7kB / 126B
readyz	4	interviewly-web-1	0.0%	66.27MiB / 7.75GiB	0.8%	15.7MB / 521MB
readyz	4	interviewly-kibana-1	8.8%	595MiB / 7.75GiB	7.5%	15.1MB / 90.9MB
readyz	4	interviewly-elasticsearch-1	3.0%	1.047GiB / 7.75GiB	13.5%	86.8MB / 14.6MB
readyz	4	interviewly-db-1	40.1%	122.5MiB / 7.75GiB	1.5%	339MB / 4.35GB
readyz	4	interviewly-cache-1	16.1%	13.7MiB / 7.75GiB	0.2%	72.9MB / 57.1MB
readyz	4	interviewly-edge-1	148.3%	36.95MiB / 7.75GiB	0.5%	2.61GB / 2.89GB
me	4	interviewly-api-4	133.3%	154.2MiB / 7.75GiB	1.9%	72.8MB / 42.2MB
me	4	interviewly-api-3	124.9%	154.7MiB / 7.75GiB	1.9%	72.8MB / 42.2MB
me	4	interviewly-api-2	137.0%	162.2MiB / 7.75GiB	2.0%	722MB / 256MB
me	4	interviewly-worker-1	0.0%	51.15MiB / 7.75GiB	0.6%	295kB / 736kB
me	4	interviewly-api-1	108.5%	185.6MiB / 7.75GiB	2.3%	3.83GB / 1.31GB
me	4	interviewly-bucket-1	0.0%	90.89MiB / 7.75GiB	1.1%	44.1kB / 10.8kB
me	4	interviewly-mail-1	0.0%	7.785MiB / 7.75GiB	0.1%	10.7kB / 126B
me	4	interviewly-web-1	0.0%	67.87MiB / 7.75GiB	0.9%	15.7MB / 521MB
me	4	interviewly-kibana-1	7.4%	595MiB / 7.75GiB	7.5%	15.1MB / 91.1MB
me	4	interviewly-elasticsearch-1	3.2%	1.047GiB / 7.75GiB	13.5%	87MB / 14.6MB
me	4	interviewly-db-1	67.4%	128.5MiB / 7.75GiB	1.6%	361MB / 4.5GB
me	4	interviewly-cache-1	0.5%	13.68MiB / 7.75GiB	0.2%	74.3MB / 58.2MB
me	4	interviewly-edge-1	66.2%	34.77MiB / 7.75GiB	0.4%	2.66GB / 2.95GB
my-interviews	4	interviewly-api-4	132.3%	158.7MiB / 7.75GiB	2.0%	256MB / 65.9MB
my-interviews	4	interviewly-api-3	137.3%	159MiB / 7.75GiB	2.0%	255MB / 65.7MB
my-interviews	4	interviewly-api-2	131.6%	163.1MiB / 7.75GiB	2.1%	904MB / 280MB
my-interviews	4	interviewly-worker-1	0.0%	51.15MiB / 7.75GiB	0.6%	299kB / 749kB
my-interviews	4	interviewly-api-1	138.5%	188.3MiB / 7.75GiB	2.4%	4.01GB / 1.34GB
my-interviews	4	interviewly-bucket-1	0.1%	90.89MiB / 7.75GiB	1.1%	44.1kB / 10.8kB
my-interviews	4	interviewly-mail-1	0.0%	7.797MiB / 7.75GiB	0.1%	10.7kB / 126B
my-interviews	4	interviewly-web-1	0.0%	69.42MiB / 7.75GiB	0.9%	15.7MB / 521MB
my-interviews	4	interviewly-kibana-1	7.0%	595MiB / 7.75GiB	7.5%	15.2MB / 91.4MB
my-interviews	4	interviewly-elasticsearch-1	0.5%	1.047GiB / 7.75GiB	13.5%	87.2MB / 14.6MB
my-interviews	4	interviewly-db-1	78.1%	134.2MiB / 7.75GiB	1.7%	385MB / 5.22GB
my-interviews	4	interviewly-cache-1	0.8%	13.68MiB / 7.75GiB	0.2%	74.3MB / 58.2MB
my-interviews	4	interviewly-edge-1	40.0%	33.27MiB / 7.75GiB	0.4%	2.73GB / 3.03GB

scenario	replicas	container	cpu	memory	mem %	net io
interview-state	4	interviewly-api-4	153.6%	161.8MiB / 7.75GiB	2.0%	345MB / 109MB
interview-state	4	interviewly-api-3	115.3%	173.4MiB / 7.75GiB	2.2%	346MB / 109MB
interview-state	4	interviewly-api-2	162.4%	172.7MiB / 7.75GiB	2.2%	994MB / 323MB
interview-state	4	interviewly-worker-1	0.2%	51.15MiB / 7.75GiB	0.6%	303kB / 766kB
interview-state	4	interviewly-api-1	114.4%	188.2MiB / 7.75GiB	2.4%	4.1GB / 1.38GB
interview-state	4	interviewly-bucket-1	0.1%	90.91MiB / 7.75GiB	1.1%	44.1kB / 10.8kB
interview-state	4	interviewly-mail-1	0.0%	7.797MiB / 7.75GiB	0.1%	10.7kB / 126B
interview-state	4	interviewly-web-1	0.0%	69.62MiB / 7.75GiB	0.9%	15.7MB / 521MB
interview-state	4	interviewly-kibana-1	4.5%	595MiB / 7.75GiB	7.5%	15.2MB / 91.6MB
interview-state	4	interviewly-elasticsearch-1	0.8%	1.047GiB / 7.75GiB	13.5%	87.5MB / 14.7MB
interview-state	4	interviewly-db-1	94.1%	145.5MiB / 7.75GiB	1.8%	413MB / 5.58GB
interview-state	4	interviewly-cache-1	2.5%	13.69MiB / 7.75GiB	0.2%	75.2MB / 58.6MB
interview-state	4	interviewly-edge-1	20.9%	33.23MiB / 7.75GiB	0.4%	2.88GB / 3.18GB
web-home	4	interviewly-api-4	0.2%	42.59MiB / 7.75GiB	0.5%	359MB / 118MB
web-home	4	interviewly-api-3	0.3%	42.5MiB / 7.75GiB	0.5%	359MB / 118MB
web-home	4	interviewly-api-2	0.4%	171.1MiB / 7.75GiB	2.2%	1.01GB / 332MB
web-home	4	interviewly-worker-1	0.3%	51.15MiB / 7.75GiB	0.6%	307kB / 780kB
web-home	4	interviewly-api-1	5.0%	64.67MiB / 7.75GiB	0.8%	4.12GB / 1.39GB
web-home	4	interviewly-bucket-1	0.0%	90.91MiB / 7.75GiB	1.1%	44.1kB / 10.8kB
web-home	4	interviewly-mail-1	0.0%	7.797MiB / 7.75GiB	0.1%	10.7kB / 126B
web-home	4	interviewly-web-1	165.9%	315.6MiB / 7.75GiB	4.0%	19.2MB / 636MB
web-home	4	interviewly-kibana-1	3.9%	595MiB / 7.75GiB	7.5%	15.2MB / 91.9MB
web-home	4	interviewly-elasticsearch-1	3.9%	1.047GiB / 7.75GiB	13.5%	87.7MB / 14.7MB
web-home	4	interviewly-db-1	0.0%	145.4MiB / 7.75GiB	1.8%	418MB / 5.63GB
web-home	4	interviewly-cache-1	0.5%	13.68MiB / 7.75GiB	0.2%	75.5MB / 58.7MB
web-home	4	interviewly-edge-1	14.2%	34.75MiB / 7.75GiB	0.4%	3.03GB / 3.33GB

Generator load

The generator shares the host with the stack. A run whose generator core utilisation approaches 100 per cent of one core is bounded by the generator rather than by the replicas, and its throughput is a floor rather than a capacity.

scenario	replicas	generator core %	requests	rps
healthz	1	61.9%	74,207	6400
readyz	1	58.2%	54,626	4698
me	1	26.7%	13,172	1129
my-interviews	1	15.6%	4,686	400
interview-state	1	9.4%	1,709	144
web-home	1	38.6%	1,508	128
healthz	2	69.8%	68,438	5974
readyz	2	68.0%	42,027	3685
me	2	36.3%	14,226	1247
my-interviews	2	21.3%	6,163	535
interview-state	2	19.5%	2,896	248
web-home	2	37.0%	1,522	128
healthz	4	80.1%	84,575	7462
readyz	4	80.8%	48,945	4437
me	4	46.1%	17,174	1518
my-interviews	4	27.9%	5,345	524
interview-state	4	27.8%	3,608	324

scenario	replicas	generator core %	requests	rps
web-home	4	30.6%	1,383	117

8. Findings the data supports

- healthz throughput went from 6400 rps at 1 replica(s) to 7462 rps at 4 replica(s), a factor of 1.17 (connections = 64), while api container cpu summed to 333.2% across 4 api container(s).
- readyz did not scale: 4698 rps at 1 replica(s) versus 4437 rps at 4 replica(s) is a factor of 0.94 (connections = 64), while api container cpu summed to 375.5% across 4 api container(s).
- me throughput went from 1129 rps at 1 replica(s) to 1518 rps at 4 replica(s), a factor of 1.34 (connections = 64), while api container cpu summed to 503.8% across 4 api container(s).
- my-interviews throughput went from 400 rps at 1 replica(s) to 524 rps at 4 replica(s), a factor of 1.31 (connections = 64), while api container cpu summed to 539.7% across 4 api container(s).
- interview-state throughput went from 144 rps at 1 replica(s) to 324 rps at 4 replica(s), a factor of 2.24 (connections = 64), while api container cpu summed to 545.7% across 4 api container(s).
- web-home did not scale: 128 rps at 1 replica(s) versus 117 rps at 4 replica(s) is a factor of 0.92 (connections = 64), while api container cpu summed to 5.9% across 4 api container(s).
- The widest load-balancer spread was in interview-state at 2 replicas, connections = 64: 6a5c890c took 50.5%, c2f71fa0 took 49.5% — a gap of 1.0 percentage points between busiest and quietest instance.
- Postgres connections peaked at 42 during readyz at 4 replicas, connections = 64, with a configured pool of 10 per api process.
- The highest event-loop delay p99 was 91.55 ms on instance c2f71fa0 during interview-state at 1 replicas, connections = 64.
- The largest gap between client and server p95 was 44.77 ms in interview-state at 1 replicas, connections = 64: client p95 611.43 ms against worst-instance server p95 566.66 ms on route GET /interviews/:id/state.
- Generator core utilisation peaked at 80.8% of one core during readyz at 4 replicas, connections = 64; no run reached 90 per cent, so no throughput figure here is a measurement of the generator.
- No run reported a transport failure: requests.failures is 0 in all 36 runs.

9. Caveats

- The generator is closed loop: one in-flight request per connection, so a slow response throttles the offered load. There is no coordinated-omission correction, and reported latency is the latency of requests that were actually sent.
- Generator and system under test run on the same host, competing for the same cores and memory.
- Authentication uses a single seeded user and a single session cookie, so session lookup caching is warmer than a real multi-user population would be.
- AI providers are stubbed for this run (config.aiEnabled is false), so no scenario exercises a model call or its latency.
- Only the api service was scaled. The web tier, database, cache and edge each stayed at one instance and are shared across every replica count.
- One run per cell: each replica/scenario/connection combination was measured once, so cell-to-cell differences below run-to-run noise carry no signal.
- Docker stats, redis client counts and database counters are single samples taken mid-run, not averages over the measurement window.